

The logo for the GPU Technology Conference is located in the top-left corner. It consists of a green rectangular box with a small green triangle pointing downwards from its bottom-left corner. Inside the box, the text "GPU" is written in a large, bold, white sans-serif font, and "TECHNOLOGY CONFERENCE" is written in a smaller, white sans-serif font to its right.

**GPU** TECHNOLOGY  
CONFERENCE


# GPU Direct for Video




# Agenda

- Video Processing Workflow
- Video I/O with GPU
- GPUDirect for Video
- Why Not Peer to Peer
- SDK Availability
- Conclusions

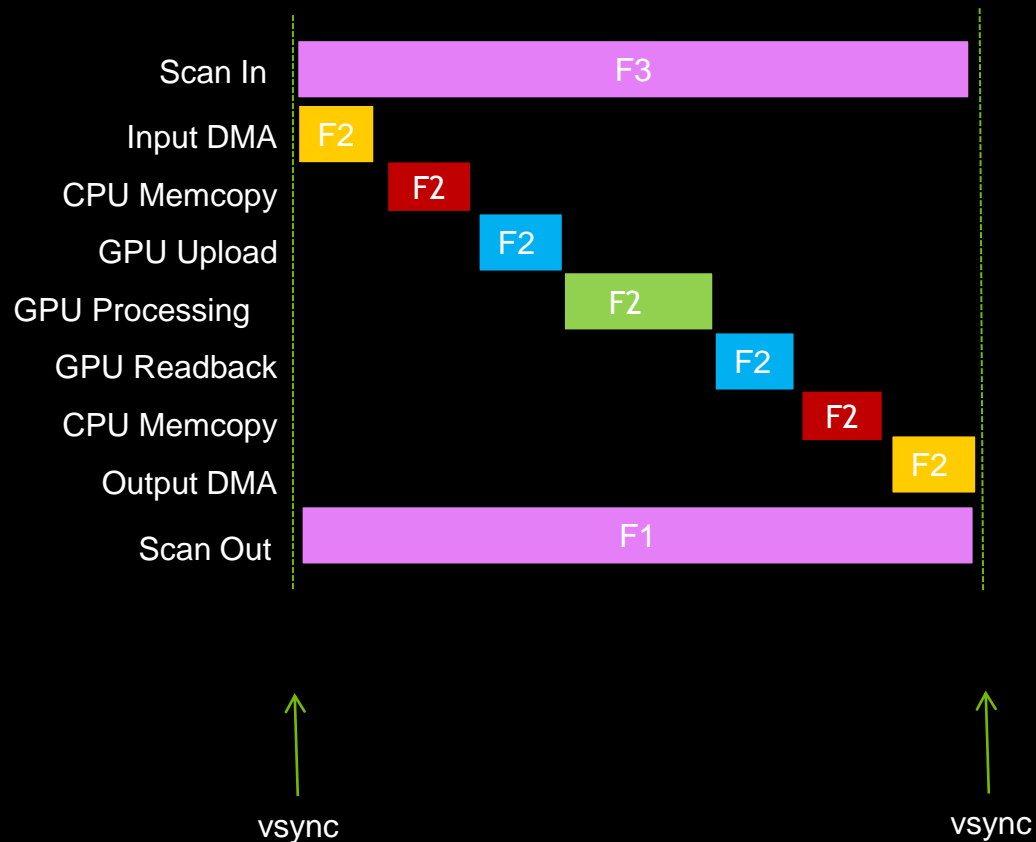
# Video Processing Workflow



# Video I/O with GPU: transfer path



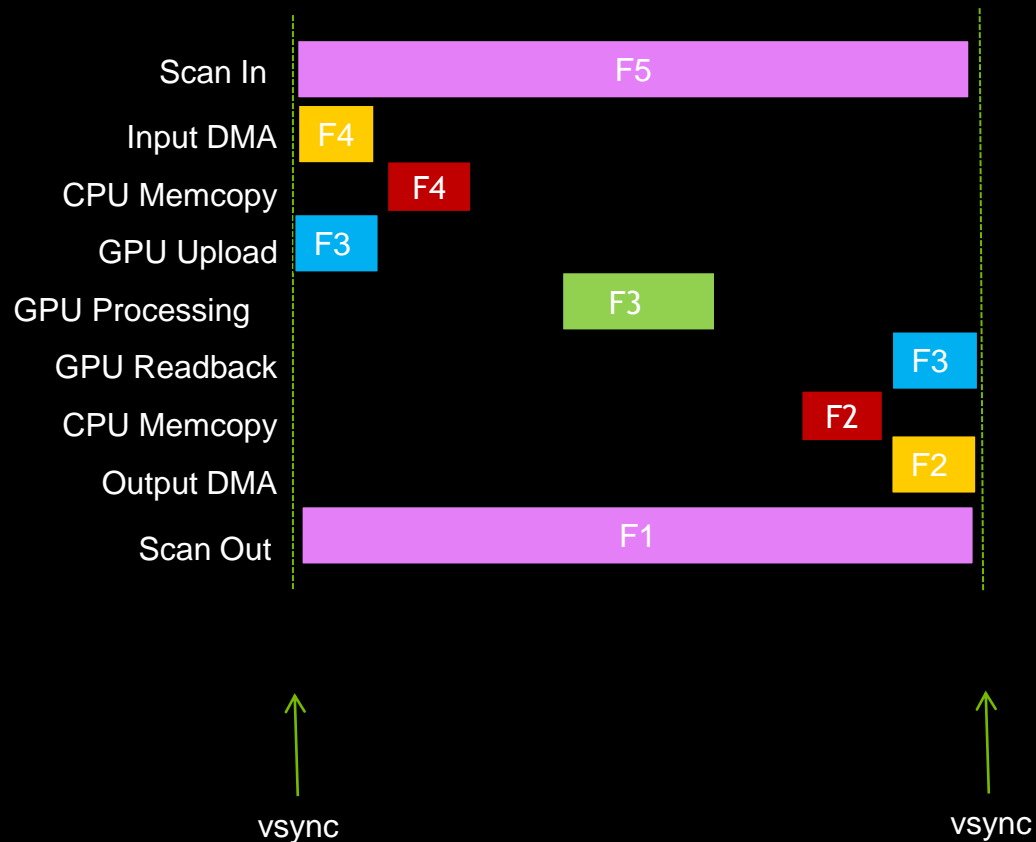
# Video I/O with GPU: inefficiencies



2 frames of latency

- I/O <-> GPU through GPU unshareable system memory buffer
- GPU synchronous transfers

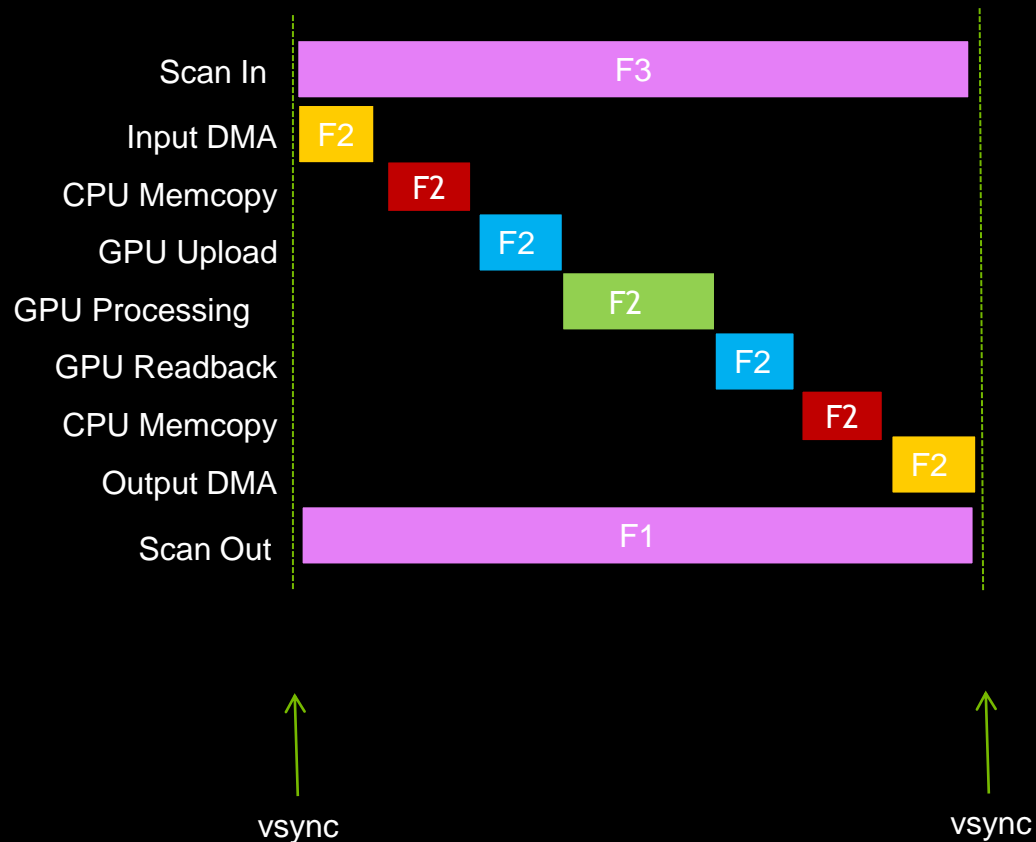
# Video I/O with GPU: improvements



■ Increasing latency

4 frames of latency!


# Video I/O with GPU: improvements



- I/O <-> GPU through GPU shareable system memory buffer

2 frames of latency

# Video I/O with GPU: improvements




2 frames of latency

- Overlapping I/O and GPU DMAs by using:
  - sub-field chunks



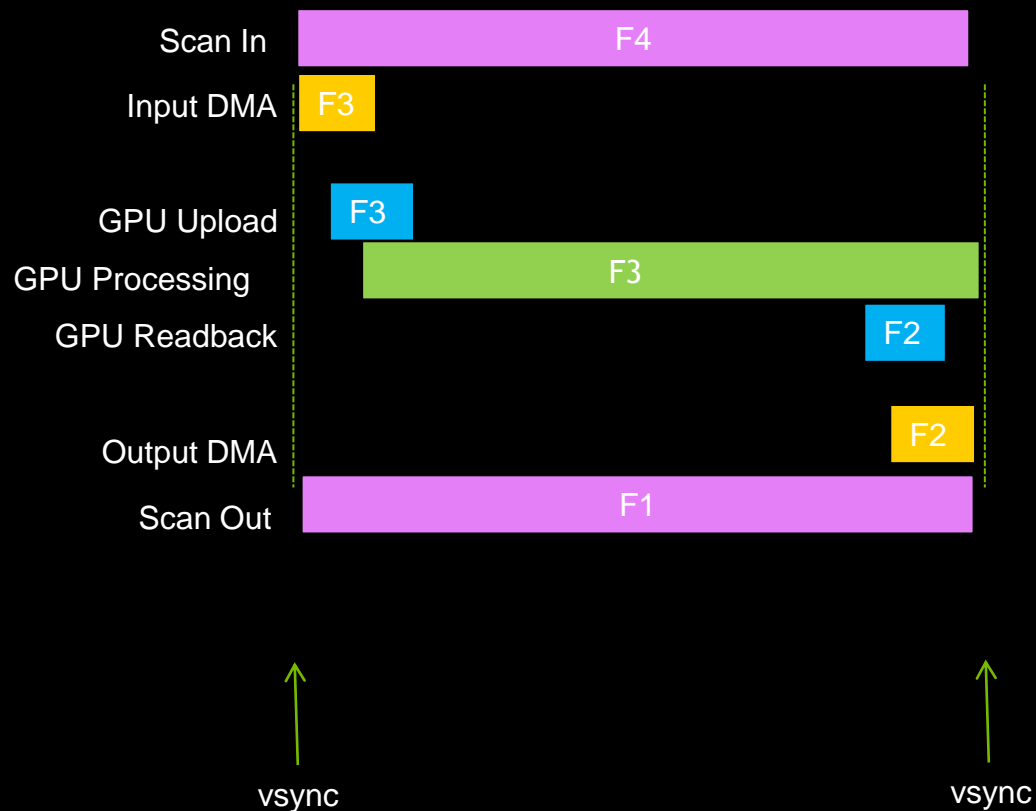
# Video I/O with GPU: improvements



2 frames of latency

- Overlapping GPU DMAs and GPU processing by using:
  - sub-field chunks
  - GPU asynchronous transfers


# Video I/O with GPU: improvements



- Overlapping GPU DMAs and GPU processing by using:
  - GPU asynchronous transfers
  - Increasing latency

3 frames of latency


# Video I/O with GPU: improvements



- Overlapping GPU DMAs and GPU processing by using:
  - GPU asynchronous transfers
  - Increasing latency even more

4 frames of latency = processing can occupy the entire frame!

# GPUDirect for Video



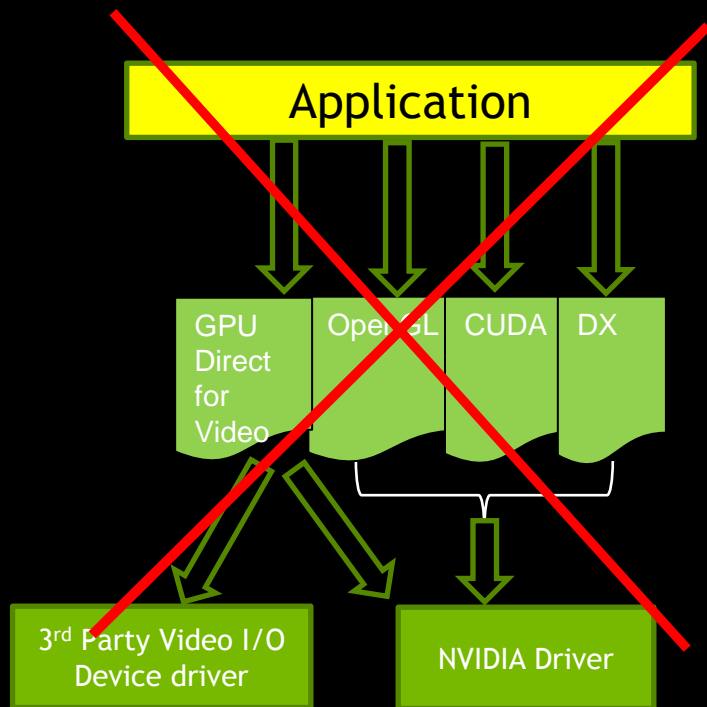
2 frames of latency!

- Unified data transfer API for all Graphics and Compute APIs' objects
  - Video oriented
  - Efficient synchronization
  - GPU shareable system memory
  - Sub-field transfers
  - GPU Asynchronous transfers

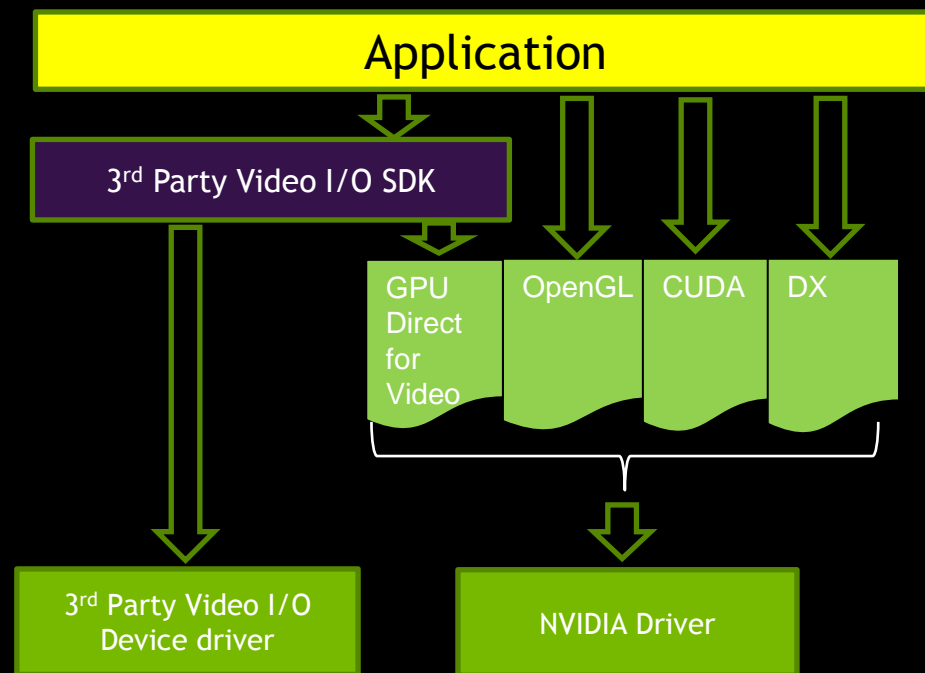


# GPUDirect for Video: Application usage

- Not this



- But this:



# GPUDirect for Video: Application Usage

Use the SDK Provided by Your Preferred Video I/O Vendor

Blackmagicdesign



# GPUDirect for Video: Application Usage

## Video Capture to OpenGL Texture

```
main()
{
    ....
    GLuint glTex;
    glGenTextures(1, &glTex);  \\ Create OpenGL texture object
    glBindTexture(GL_TEXTURE_2D, glTex);
    glTexImage2D(GL_TEXTURE_2D, 0, GL_RGB, bufferWidth, bufferHeight, 0, 0, 0, 0);
    glBindTexture(GL_TEXTURE_2D, 0);
    EXTRegisterGPUDirectTextureGL(glTexIn);  \\ Register texture with 3rd party Video I/O SDK

    while(!quit)
    {
        EXTBegin(glTexIn);  \\ Release texture from Video I/O SDK
        Render(glTexIn);  \\ Use the texture
        EXTEnd(glTexIn);  \\ Release texture back to Video I/O SDK
    }
    EXTUnregisterGPUDirectTextureGL(glTexIn);  \\ Unregister texture with 3rd party Video I/O SDK
}
```

# Results

Optimal transfer time for 4-component 8-bit 1080p video:

$$\text{transfer time} \cong \frac{\text{frame size}}{\text{PCIe bandwidth}} * 2 \cong \frac{497664000 \text{ bytes per second}}{60000000000^* \text{ bytes per second}} * \frac{1}{60} * 2$$

$$\text{transfer time} \cong 2.397 \text{ msec}$$

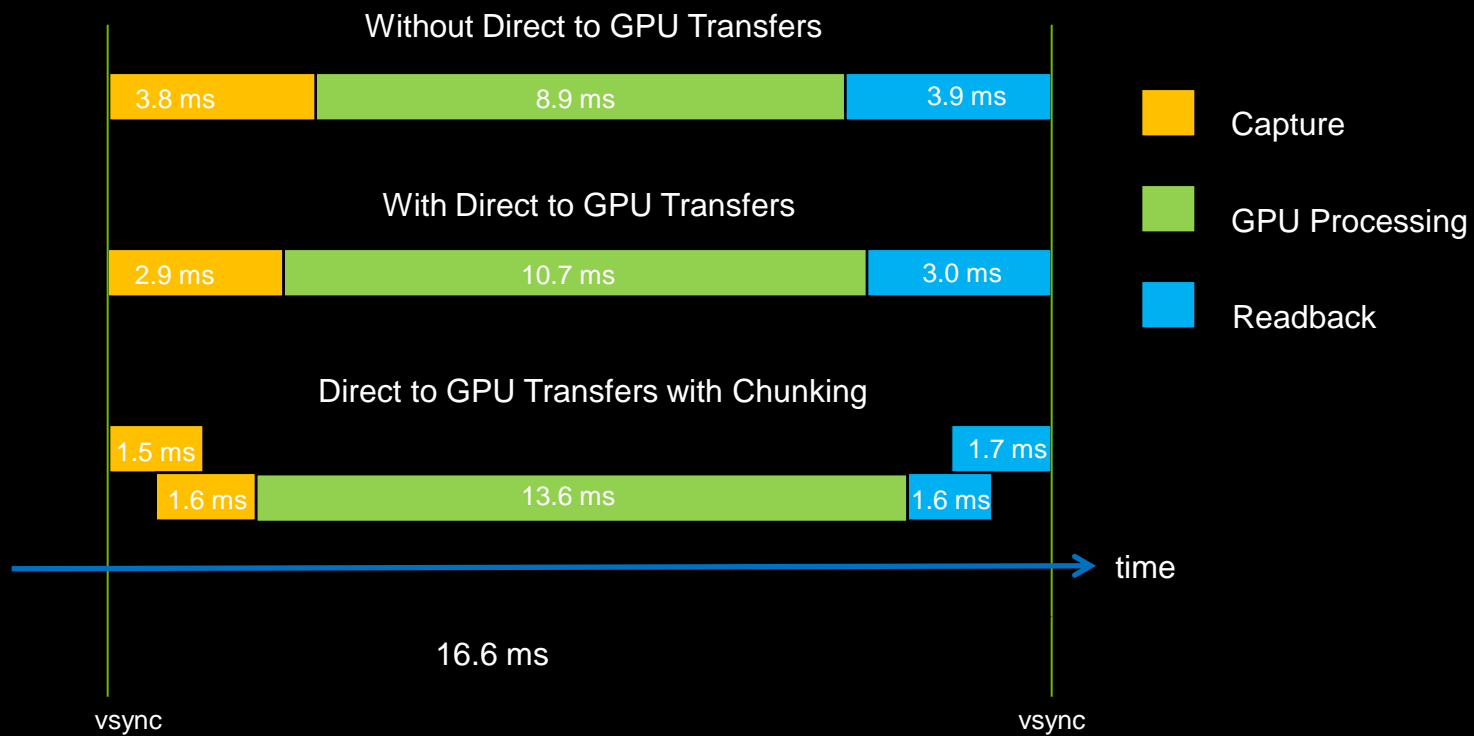
\* Although Gen 2 PCI Express bandwidth is specified at 8.0GB / sec, the maximum achievable is ~6.0 GB/sec

Direct to GPU video transfer time for 4-component 8-bit 1080p video:

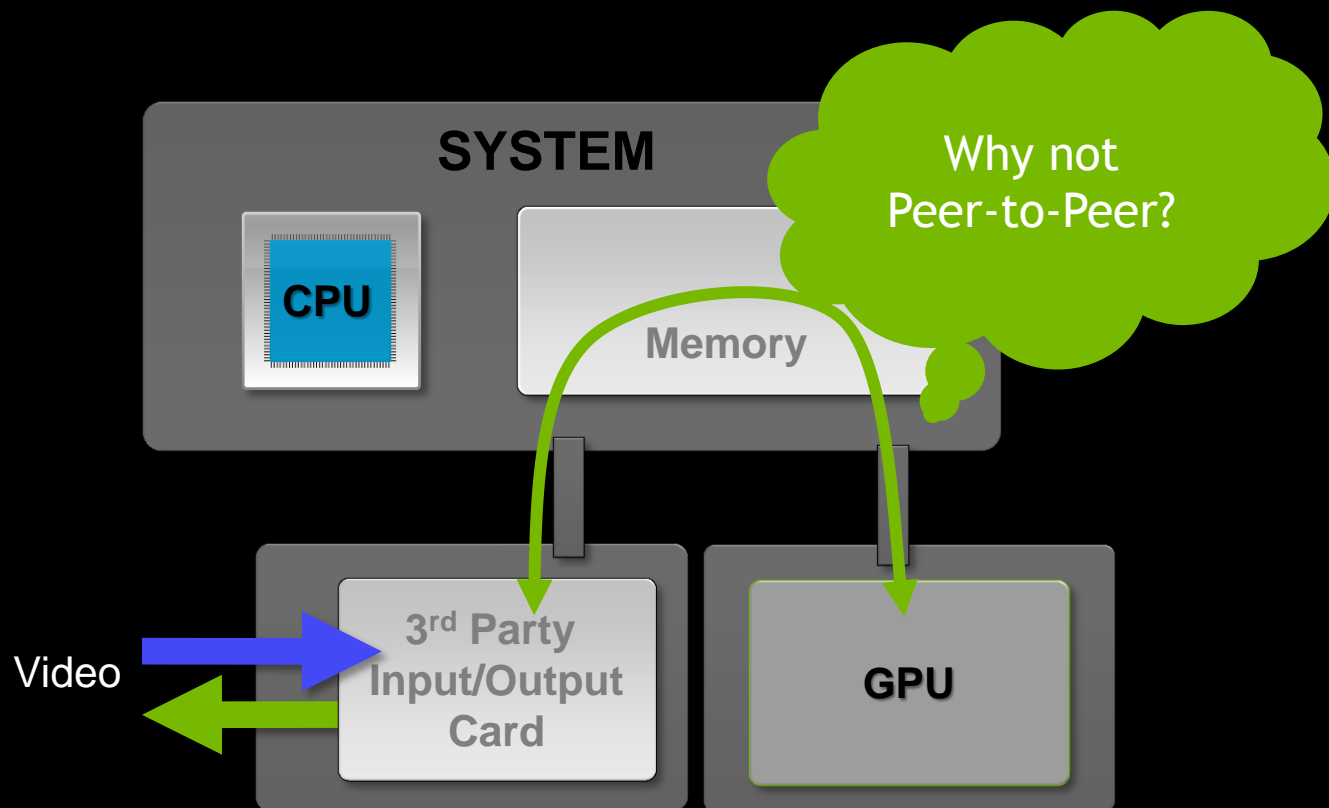
Capture Latency		Playout Latency	
ExtDev to SysMem	SysMem to GPU	GPU to SysMem	SysMem to ExtDev
1.4ms	1.5ms	1.5ms	1.5ms
2.9 ms		3.0ms	



# Results

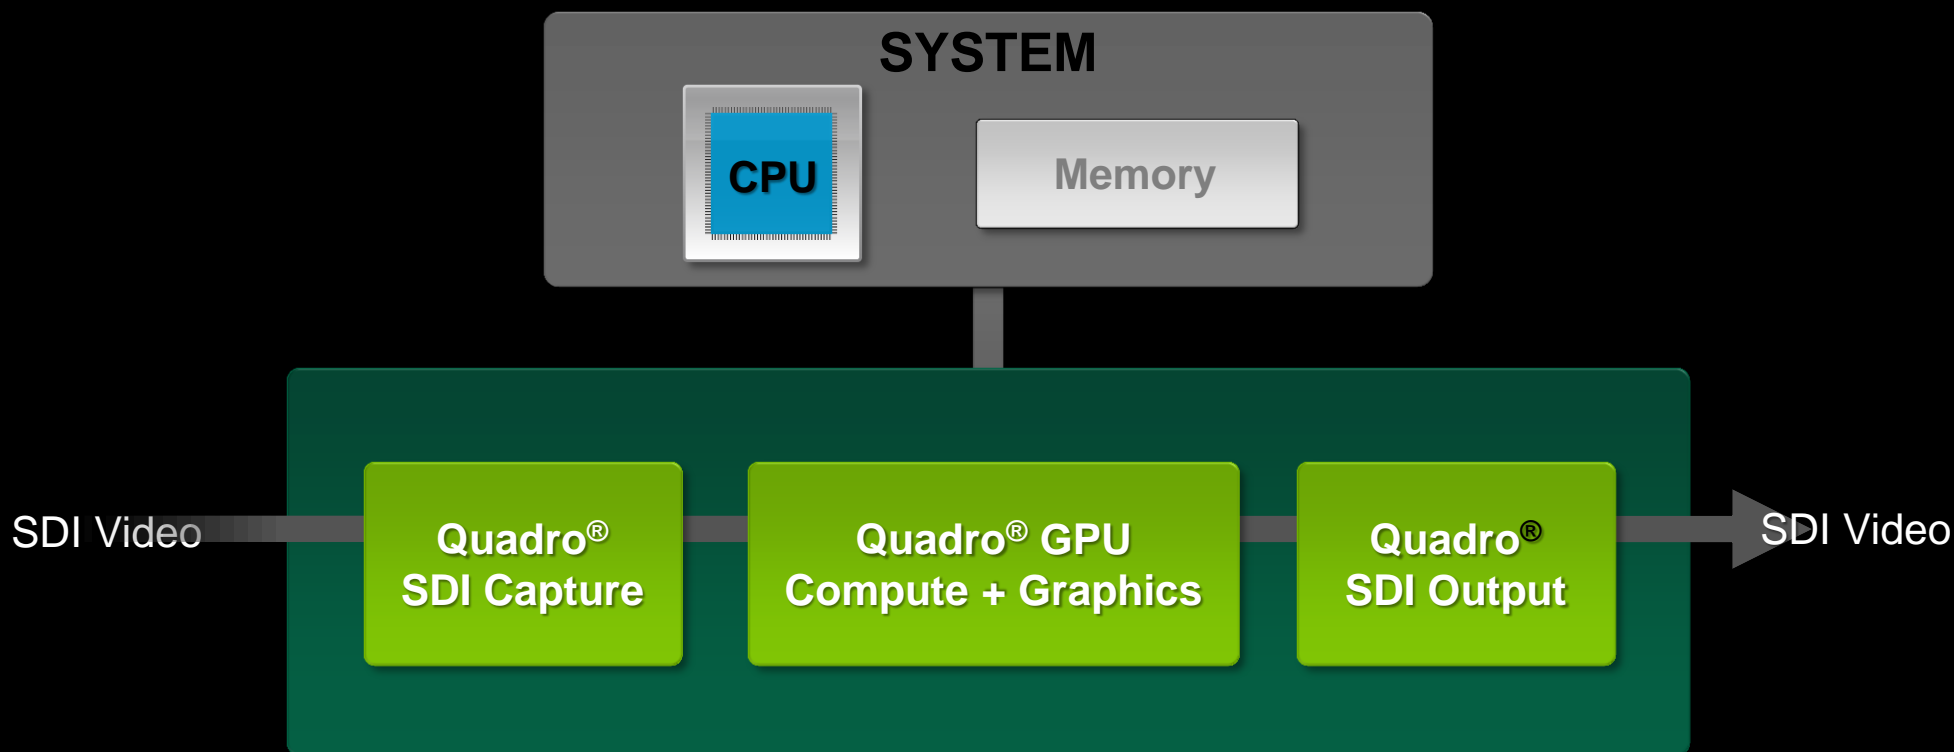


# Video I/O with GPU: transfer path

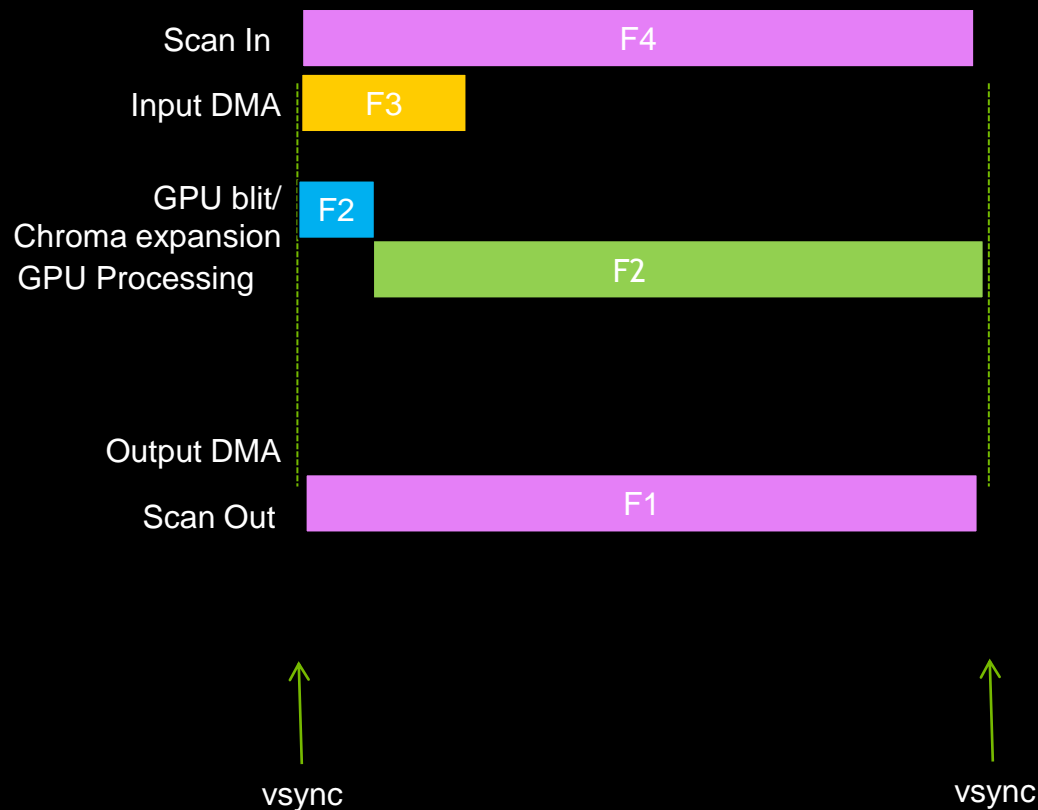


# NVIDIA Digital Video Pipeline: Peer-to-Peer Communication

*NVIDIA SDI capture and output cards only*



# NVIDIA Digital Video Pipeline



3 frames of latency!

- The Input DMA is peer-to-peer of 4 inputs
- The output is a direct scanout from the GPU

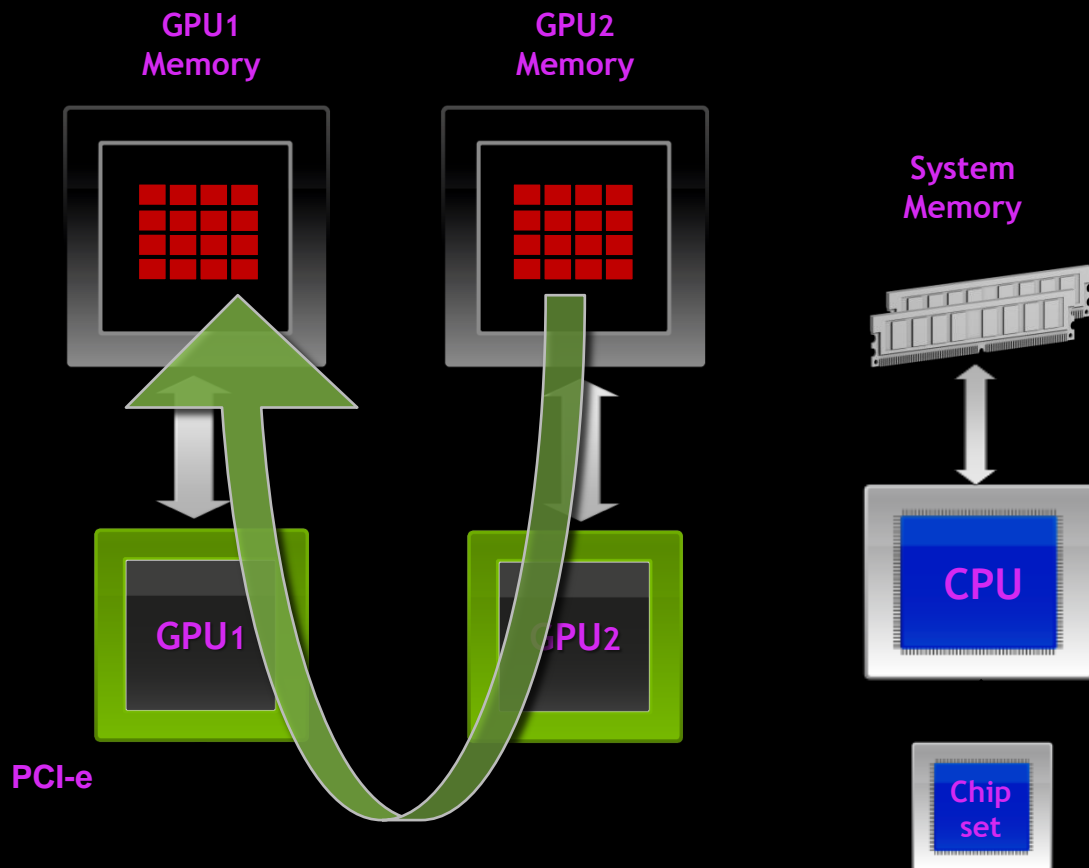
**BUT...**

- Fixed 3 frames of latency
- Each I/O board must be tied to a single GPU




# NVIDIA GPUDirect™ v2.0: Peer-to-Peer Communication

*Direct Transfers between GPUs with CUDA*



**Only one copy required:**  
1. `cudaMemcpy(GPU2, GPU1)`

# NVIDIA GPUDirect™ now supports RDMA



Server 1

Server 2

*Linux only, HPC centric*

# Why Not Peer-to-Peer?

- Supportability
- Same performance
- IOH-to-IOH communication issues
- Limited PCIe slot options due to lane allocations
- Support Graphics APIs as well as CUDA
- Multi-GPU Support!
- Multi-OS support!

## Video I/O Card Vendors

Use the NVIDIA GPU Direct for Video SDK:

<http://developer.nvidia.com/nvidia-gpudirect%E2%84%A2-video>

- Samples (OpenGL, D3D9, D3D11, CUDA)
- Programming Guide
- Windows7, Linux
- Static and Shared Libraries



# Conclusions

- Lowest latency video I/O in and out of the GPU
- Optimal transfer times
- Optimal GPU processing time
- Supports OpenGL and DirectX as well as CUDA
- Does not require sophisticated programming.
- Scales to multiple GPUs